

Monitoring and tuning Apache Spark

TANG Gen

Outline

- Components of execution
 - Jobs, stages and tasks
 - Evaluation of DAGs
 - Spark UI
- Spark UI
- Spark logs
- Key performance considerations

/Jobs, stages and tasks

- How does a user program get translated into units of physical execution: jobs, stages and tasks ?

[illegible]

The screenshot shows the Spark UI interface in a web browser. The top navigation bar includes links for Jobs, Stages, Storage, Environment, and Executors. The main heading is "Spark Jobs (?)". Below this, summary statistics are provided: Total Duration: 2.9 min, Scheduling Mode: FAIR, Active Jobs: 1, Completed Jobs: 19, and Failed Jobs: 15. The "Active Jobs (1)" section is currently empty. The "Completed Jobs (19)" section displays a table with 19 rows, each representing a completed job. Each row includes the Job ID, Description, Submitted time, Duration, Stages (Succeeded/Total), and Tasks (Succeeded/Total). The tasks are represented by blue progress bars. The bottom of the screenshot shows a status message: "Waiting for localhost..."

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
34	Job with delays count at UIWorkloadGenerator.scala:85	2014/11/27 13:30:24	1 s	0/1	10/100

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
30	Single Shuffle count at UIWorkloadGenerator.scala:64	2014/11/27 13:30:04	0.1 s	1/1 (1 skipped)	100/100 (100 skipped)
29	Cache and Count count at UIWorkloadGenerator.scala:63	2014/11/27 13:29:59	0.2 s	1/1	100/100
27	Job with delays count at UIWorkloadGenerator.scala:85	2014/11/27 13:29:49	10 s	1/1	100/100
28	Count count at UIWorkloadGenerator.scala:62	2014/11/27 13:29:54	0.1 s	1/1	100/100
23	Single Shuffle count at UIWorkloadGenerator.scala:64	2014/11/27 13:29:29	0.1 s	1/1 (1 skipped)	100/100 (100 skipped)
22	Cache and Count count at UIWorkloadGenerator.scala:63	2014/11/27 13:29:24	0.3 s	1/1	100/100
20	Job with delays count at UIWorkloadGenerator.scala:85	2014/11/27 13:29:14	10 s	1/1	100/100
21	Count count at UIWorkloadGenerator.scala:62	2014/11/27 13:29:19	0.2 s	1/1	100/100

Waiting for localhost...

Components of execution

/Jobs, stages and tasks

- Consider a very simple example
 - Read a log file
 - Split into words and remove empty lines
 - Extract the log level and do a count

```
val input = sc.textFile("log.txt")

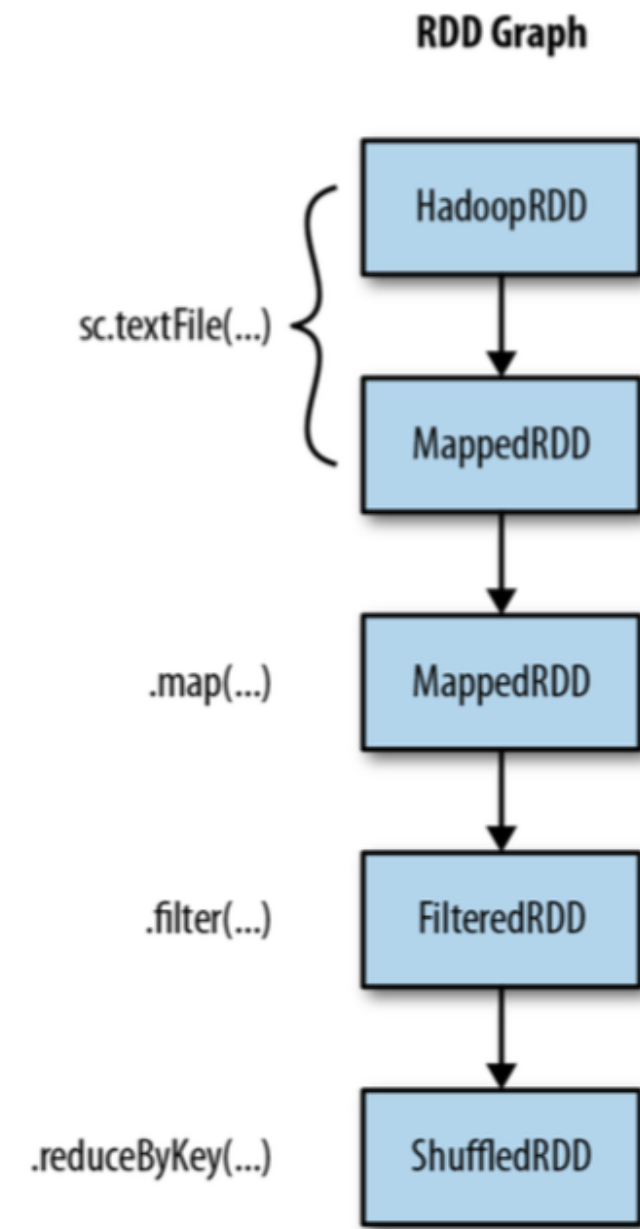
val tokenized = input
  .filter(line => line.size > 0)
  .map(line => line.split(" "))

val counts = tokenized
  .map(words => (words(0), 1))
  .reduceByKey({case (a, b) => a + b})
```

Components of execution

/Jobs, stages and tasks

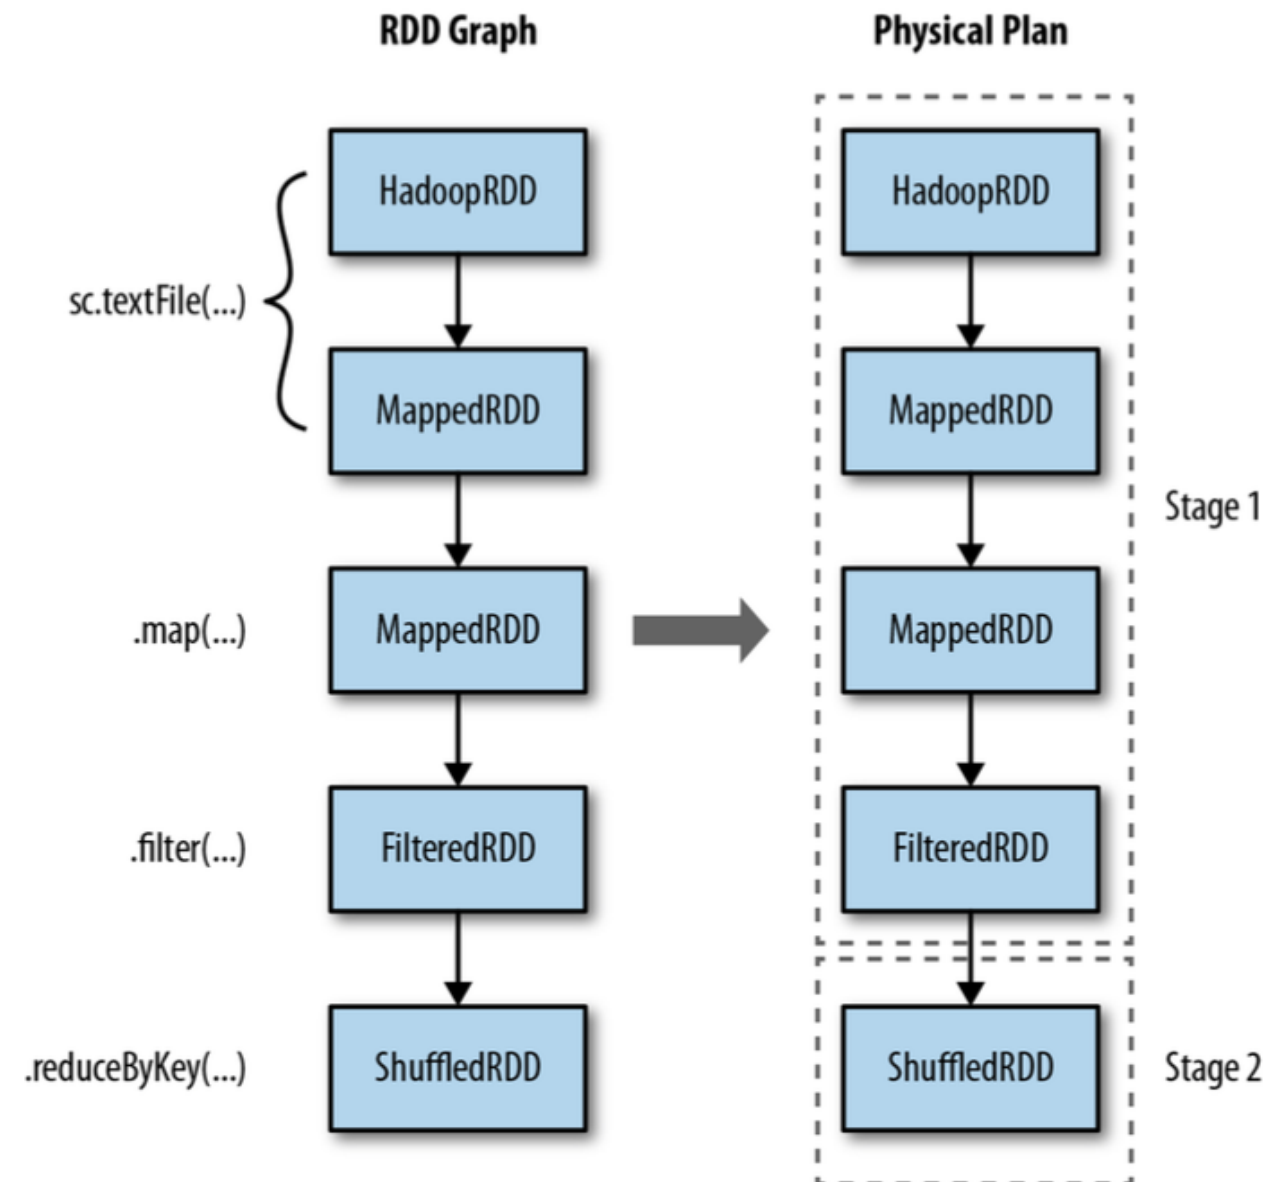
- If we execute the above commands in Spark shell, the program has not performed any actions
- Transformations build up a **DAG**, but don't "do anything"
- DAG is directed acyclic graph which will remember the lineage of RDD, including parents, dependencies, etc. It provides fault-tolerance of Spark
- Spark's scheduler creates a physical execution plan according to DAG when an action runs on RDD



Components of execution

/Jobs, stages and tasks

- If we add an action to RDD counts, for example, `counts.count()`
- Spark's scheduler starts at the final RDD being computed and works backward to find what it much compute
- The scheduler outputs a computation stages for DAG. Each stage has tasks for each partition in RDD. Those stages are executed in reverse order



Components of execution

/Evaluation of DAGs

- Before, we say that “actions” force the evaluation of RDD which is True. But Let’s forget it for a moment
- However, in fact, DAGs are materialized through a method `sc.runJob` in Spark Core

```
def runJob[T, U](  
  rdd: RDD[T],  
  partitions: Seq[Int],  
  func: (Iterator[T] => U)  
): Array[U]
```

RDD to compute

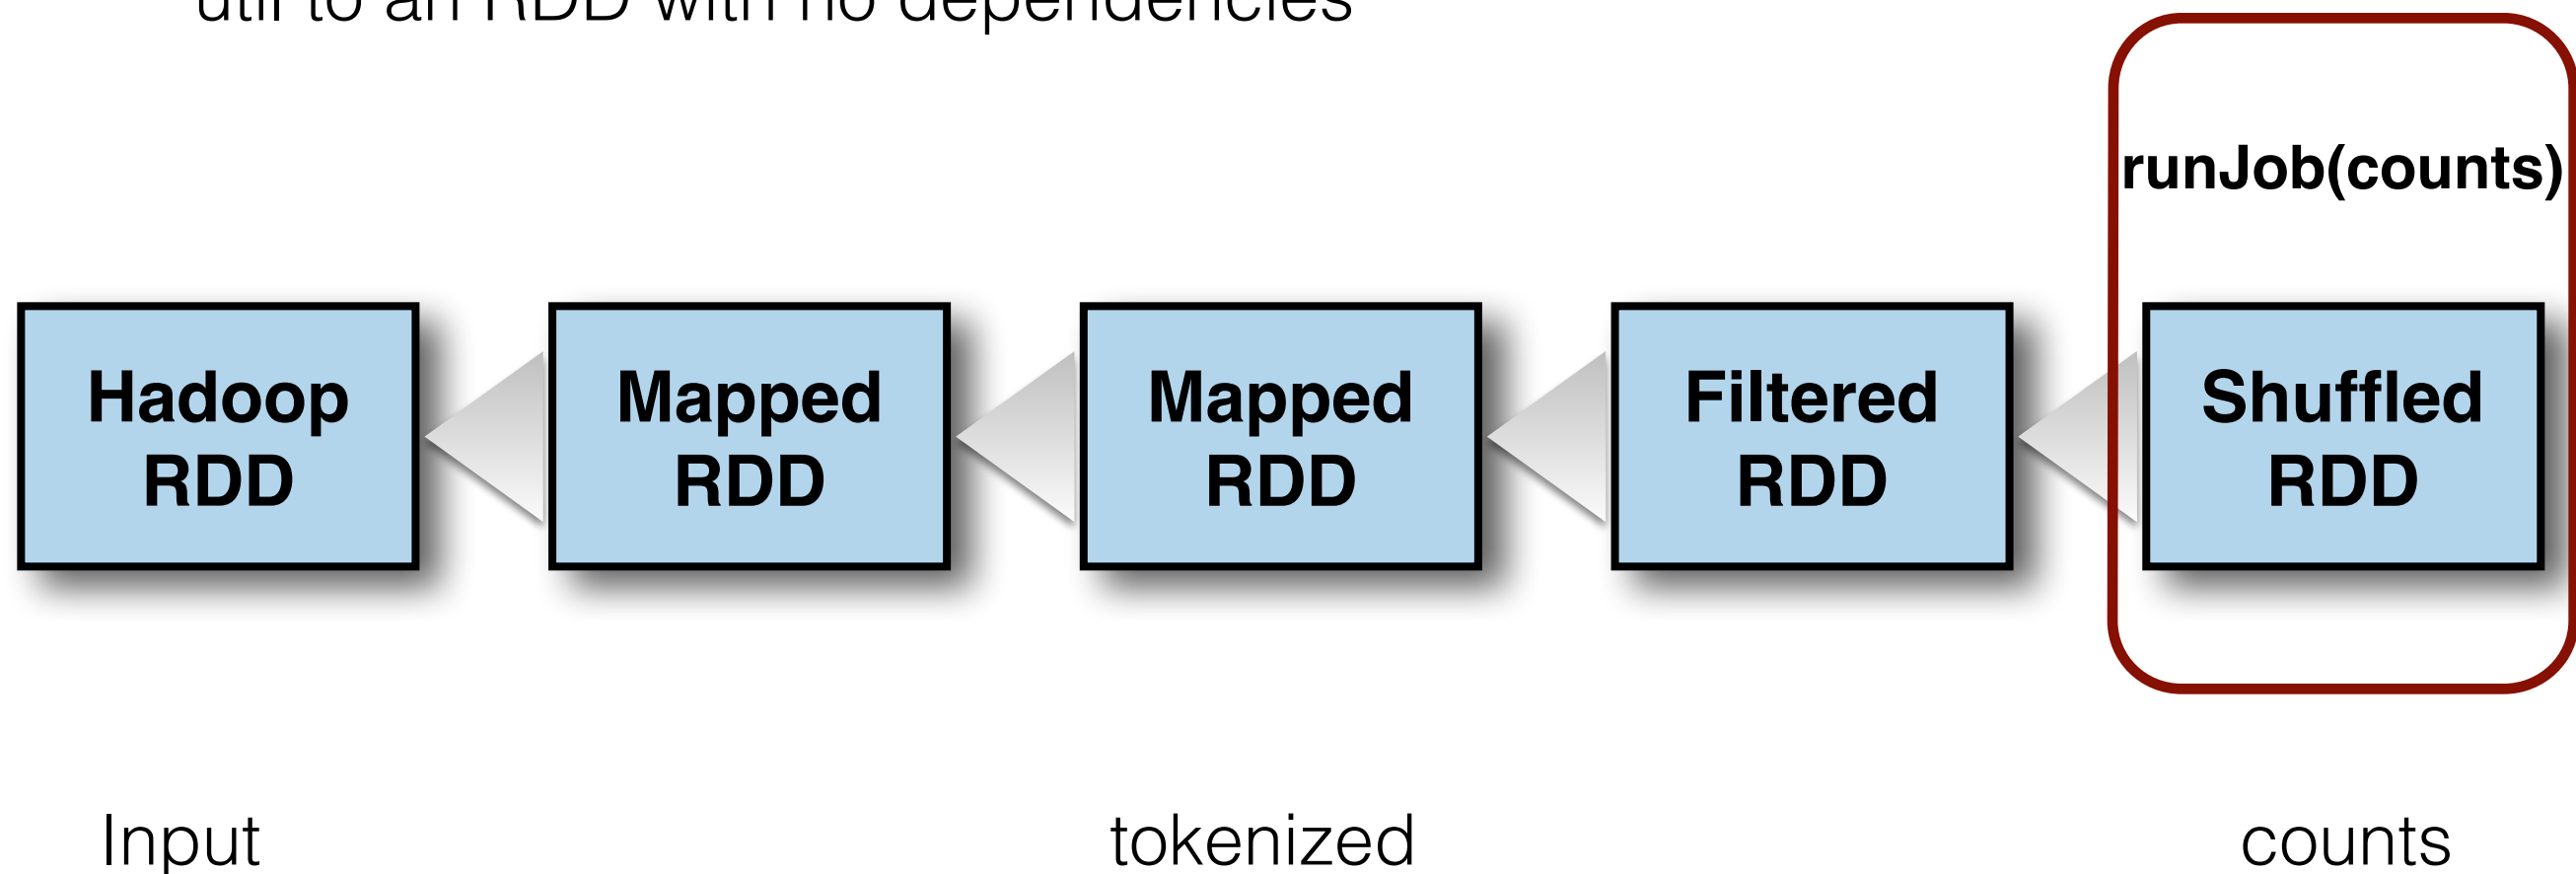
Which partitions

Function to produce
results

Components of execution

/Evaluation of DAGs

- RunJob needs to compute RDD parents, parents' parents and etc, until to an RDD with no dependencies



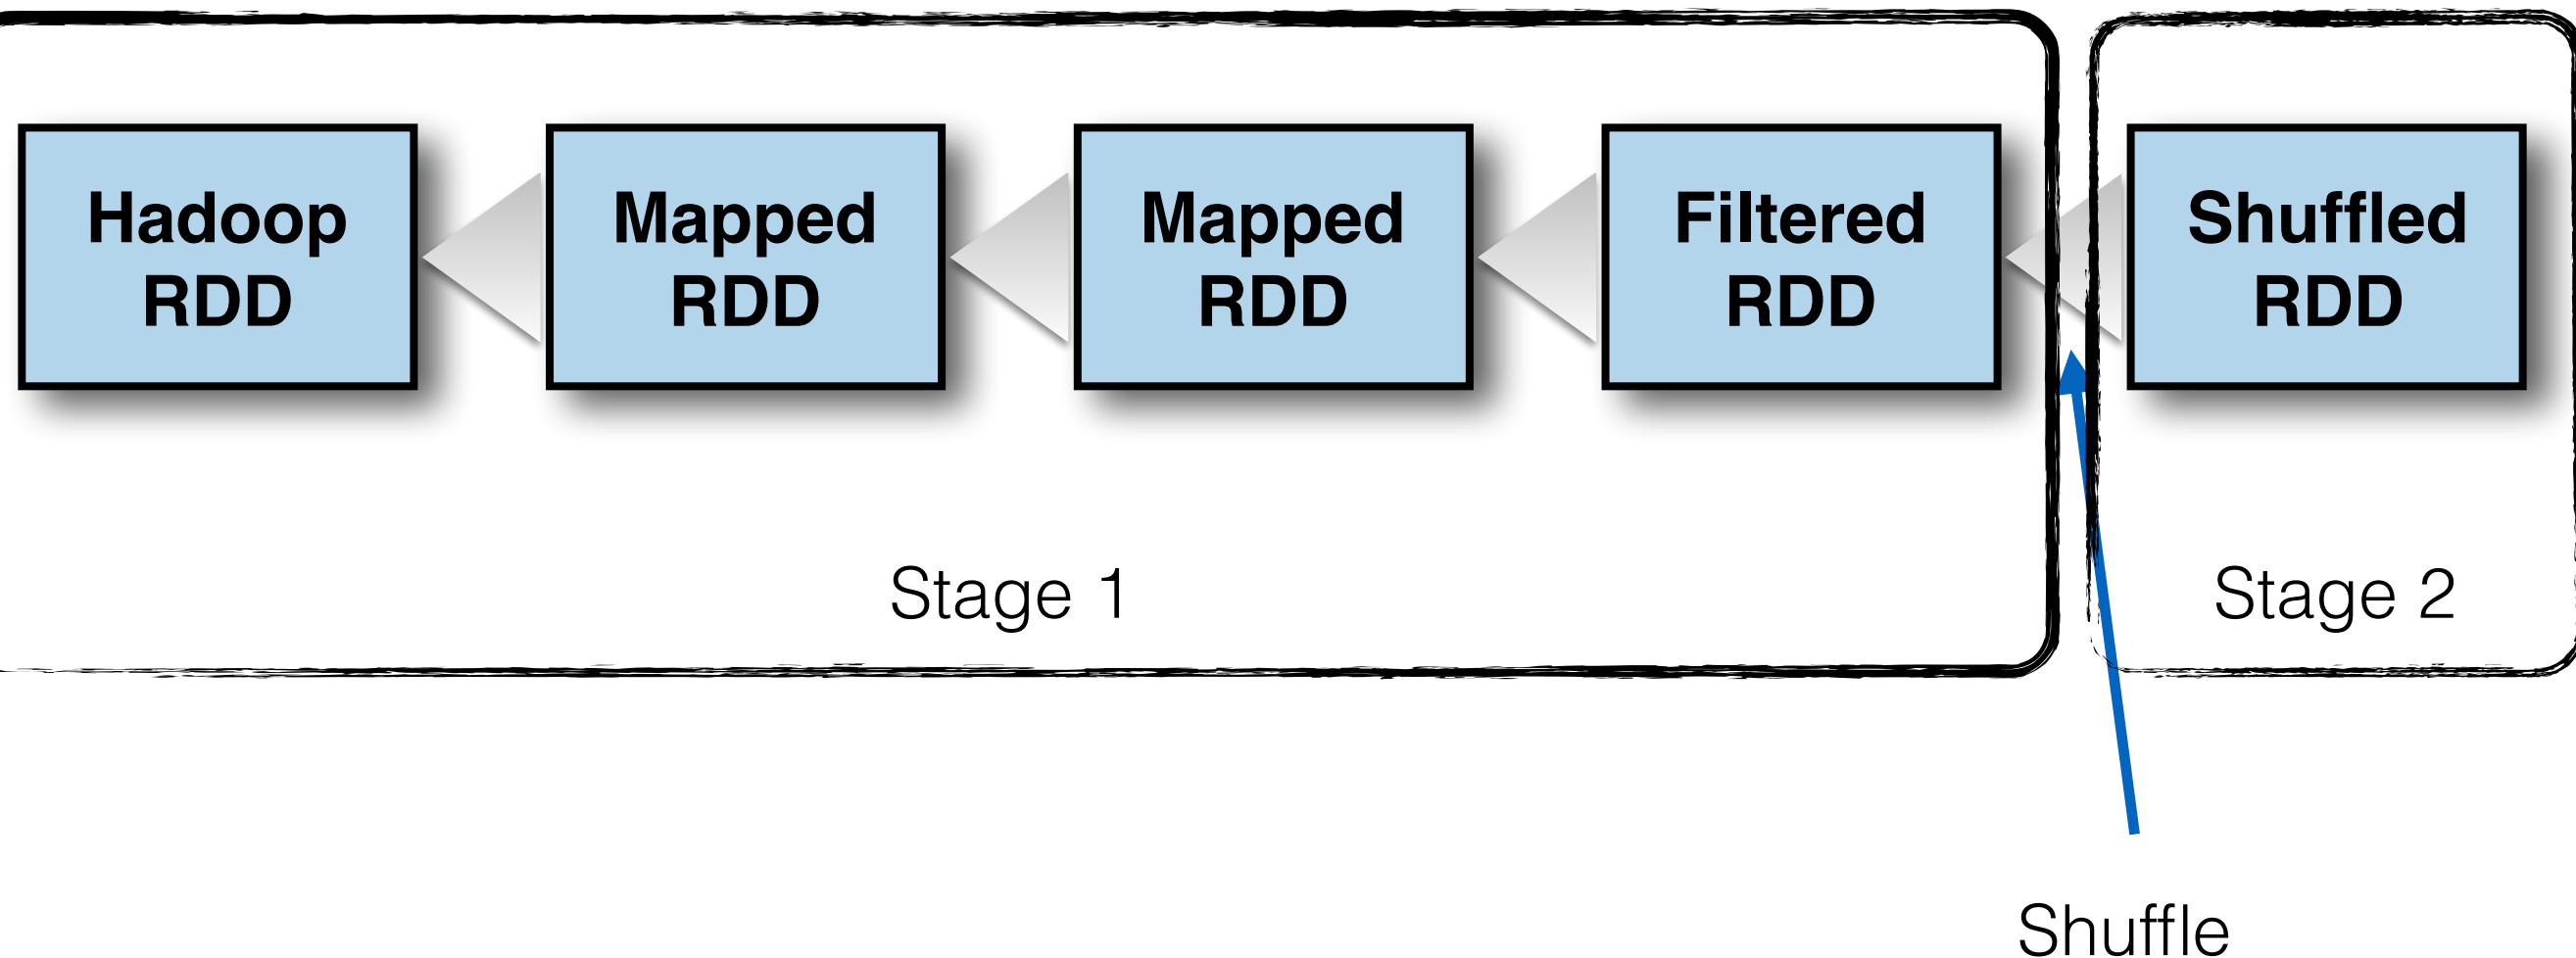
Components of execution

/Physical optimization

- Certain types of transformation can be **pipelined**
- If dependent RDD's have already been cached (or persisted in a shuffle), the graph can be **truncated**
- Once pipelining and truncation occur, Spark produces a set of **stages**, each stage is composed of **tasks**

Components of execution

/Physical optimization




Components of execution

/Stage graph

- In stage 1, each task will

1. Read hadoop file
2. Perform maps and filters
3. Write partial sums



- In stage 2, each task will

1. Read partial sums
2. Invoke user function passed to runJob

Components of execution

/Count() action

```
class RDD {  
  def count(): Long = {  
    results = sc.runJob(  
      this,  
      0 until partitions.size  
      it => it.size()  
    )  
    return results.sum  
  }  
}
```


RDD = self

Partitions = all partitions

Function = size of all
partition

Components of execution

/Spark UI



Details for Job 0

Status: SUCCEEDED

Completed Stages: 2

Completed Stages (2)

Stage Id	Description		Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
1	count at <console>:28	+details	2015/02/19 21:35:56	38 ms	2/2				
0	map at <console>:25	+details	2015/02/19 21:35:56	0.1 s	2/2	72.0 B			354.0 B

Named after action calling runJob

Named after last RDD in pipeline

Components of execution

/Summary

- Definition:
 - **Jobs:** Work required to compute RDD in runJob.
 - **Stages:** A wave of work within a job, corresponding to one or more pipelined RDD's.
 - **Tasks:** A unit of work within a stage, corresponding to one RDD partition.
 - **Shuffle:** The transfer of data between stages.
- Phases occur during Spark execution:
 - User code defines a DAG of RDDs
 - Actions force translation of the DAG to an execution plan
 - Tasks are scheduled and executed on a cluster

Spark UI

/Demo time

- If Spark runs on standalone or Mesos cluster, Spark's built-in web UI is available on the machine where driver is running at port 4040 by default
- If Spark runs on YARN cluster mode, the driver runs inside the cluster, the UI is accessed through the YARN ResourceManager

Spark logs

/Driver and executor logs

- The location of Spark's log files depends on the deployment mode:
 - Standalone, application logs are directly displayed in the stand-alone master's web UI. They are stored by default in the *work/* directory of the Spark distribution on each worker.
 - In Mesos, logs are stored in the *work/* directory of a Mesos slave, and accessible from Mesos master UI
 - In YARN mode, when the application finishes, use the following command to produce a report containing logs
yarn logs -applicationID <app ID>
- In YARN mode, when the application is running, we can access the logs of certain containers via ResourceManager UI

Spark logs

/Logging system

- Spark's logging subsystem is based on *log4j*, a widely used Java library
- An example log4j configuration file is bundled with Spark at *conf/log4j.properties.template*. To customize Spark's logging,
 1. Copy the example to a file called *log4j.properties*.
 2. Modify behavior such as the root logging level (the threshold level for logging output).
 3. Add the *log4j.properties* file using the *--files* flag of *spark-submit*

Key performance considerations

/Level of Parallelism

- RDD's partitions size decides the level of parallelism during execution, which can be modified by *coalesce(Num)* or *repartition(Num)*
- Demo time
 - Too little parallelism: idle cores
 - Too much parallelism: too much shuffle

Key performance considerations

/Serialization format

- When Spark is transferring data over the network or spilling data to disk, it needs to serialize objects into a binary format
- By default, Spark will use Java's built-in serializer.
- However, Spark also support the use of Kyro, a third-party serialization library (Kyro is used in mllib library)
- To use Kyro serializer

```
val conf = SparkConf()  
conf.set("spark.serializer", "org.apache.spark.serializer.KyroSerializer")  
conf.set("spark.kyro.registrationRequired", "true")  
conf.registerKyroClasses(Array(classOf[Myclass], classOf[Myotherclass]))
```

Key performance considerations

/Memory management

- In general, memory is used for the following three ways:
 - RDD storage
 - Shuffle and aggregation buffers
 - User code
- The more detailed presentation of memory management will be in tomorrow